
Symbolic Processing for Trajectory Prediction

Alec J. Hoyland

Center for Systems Neuroscience
Boston University
Boston, MA 02215
ahoyland@bu.edu

Abstract

The abstract paragraph should be indented ½ inch (3 picas) on both the left- and right-hand margins. Use 10 point type, with a vertical spacing (leading) of 11 points. The word **Abstract** must be centered, bold, and in point size 12. Two line spaces precede the abstract. The abstract must be limited to one paragraph.

1 Introduction

Symbolic processing is an important feature of general intelligence, crucial for understanding underlying rules and latent features, to develop models of the world that can be used for prediction. In animals, simple prediction tasks such as object tracking and navigation, are necessary for survival. These tasks fall under the umbrella of symbolic processing. One of the most straightforward symbolic processing tasks which can be abstracted to a computational model is trajectory prediction.

This task involves observing a system with a state that evolves in time for a few epochs, before making predictions of the system's future state based on those observations. Real-world applications involve pedestrian and vehicle tracking, robot navigation and agility, automated physics tracking, and aeronautics. In essence, it is a type of supervised learning problem, in which the prediction model is being updated at each time step.

The goal of this paper is to put forward several theories on how this can be practically done, with an emphasis on discovering fundamental mathematical relationships which could be performed by ensembles of neurons in the brain.

1.1 What is a trajectory?

For the purpose of this section, a *trajectory* is an ordered set of states which are evolved according to one or more rules. The purpose of symbolic trajectory prediction is, if given a few states of the trajectory, to intuit the latent rules and predict future states. In a real-world scenario, the agent predicts simultaneously to the evolution of the trajectory in time, meaning that more data are available to the agent each time step.

Mathematically, a trajectory is formally represented as a sequence of values $[f^k(x)]_{k \in \mathbb{N}}$, calculated by the iterated application of a mapping f to an element x of its source. It is analogous to the time-ordered set of states in a dynamical system (e.g. a Poincaré map) (CITE) such as the function of position produced by integrating the equations of motion in Newtonian mechanics.

1.2 What is a feature?

A feature is any important defining characteristic of a snapshot of a trajectory. For example, if the trajectory represents the motion of a ball in flight across a 2-dimensional coordinate grid, the crucial features are the coordinates (x, y) parametrised by the time-coordinate t . If the color of the ball

mattered (*i.e.* isn't time-invariant), then it would also be a feature. For the purpose of this paper, the agent is assumed to be able to discern which features are important and which are not. From a biological perspective, we will take for granted the miraculous specificity of the visual system in mammals. In artificial intelligence, we will assume a sufficiently-advanced computer vision algorithm and peripherals. If, for example, three polygonal shapes were drawn on a blackboard and the agent were asked to determine the pattern, it is reasonable to neglect the minute deviations in the lines caused by the hand of an unskilled artist. Similarly, imperfections in printed symbols could be ignored as well. The best heuristic for determining what is a feature and what is not relies on comparing adjacent states of the trajectory. Major differences are features; minor differences can be neglected.

1.3 Representing a trajectory through vector transformations

One way to imagine a trajectory is to consider a series of points in the plane formed by the time-coordinate t and a function $f(t)$ which produces a vector output \vec{x} , where each element in \vec{x} is the numerical value of a feature at time t . In this situation, the trajectory is a vector parametrized by the time-coordinate. To compare adjacent states, the first time-derivative can be taken. Since time is discrete, the derivative operator is best represented by the finite difference.

A finite difference is a mathematical expression of the form $f(x + b) - f(x - a)$. If the finite difference is divided by $b - a$, the different quotient is defined. The first forward finite difference is defined by

$$Dx_n = x_n - x_{n-1} \quad (1)$$

for variable x at discrete indices n and $n - 1$. Higher-order finite differences can also be defined, such as the second-order forward finite difference:

$$D^2x_n = x_n - 2x_{n-1} + x_{n-2} \quad (2)$$

These finite differences are first-order correct in accuracy with uniform grid-spacing t_n . More accurate finite differences can be computed using more points. At minimum, approximating a numerical derivative of order k with the fewest points requires $k + 1$ points.

If two states of the trajectory x_n and x_{n-1} are known, the next point can be approximated as:

$$x_{n+1} \approx \frac{1}{\Delta} Dx_n + x_n = \frac{x_n - x_{n-1}}{\Delta} + x_n \quad (3)$$

where $\Delta = t_n - t_{n-1}$ is the difference in the time-coordinate between adjacent indices. If higher-order finite differences are used, the estimate will be more accurate, up to the derivative order equal to the order of the polynomial representation of the function. This requires knowing more past states.

Since any well-behaved function can be approximated by its (truncated) Taylor series, the discrete Taylor series generalization, given by Einar Hille, can approximate a well-behaved function using finite differences (CITE). For $t > 0$

$$f(a + t) = \lim_{\Delta \rightarrow 0^+} \sum_{n=0}^{\infty} \frac{t^n}{n! \Delta^n} D^n f(a) \quad (4)$$

This equation describes how if a function is known to be evaluable at a : $f(a)$, and finite differences can be taken, then the function can be approximated at any future point $a + t$.

The vector transformation formulation is satisfactory for predicting trajectories based on linear rules. It is mathematically equivalent to least-squares fitting, and therefore can be represented in matrix form (CITE). This results in a linear transformation called the predictor:

$$P(a) = \left[\frac{1}{n!} D^n f(a) \right]_{n \in \mathbb{N}} \quad (5)$$

The dot product of the predictor can be taken with $\left[\frac{t^n}{\Delta^n}\right]_{n \in \mathbb{N}}$ to find the predicted trajectory $f(a + t)$ for $t > 0$.

1.3.1 Example: one-dimensional trajectory

Here we consider a trajectory produced by a linear combination in a polynomial basis. The goal is to predict the one-dimensional trajectory

$$f(t) = -40 + 100t - 10t^2 \quad (6)$$

For $\Delta = 0.01$, the first and second finite differences are computed at $f(2\Delta)$. By the Hille series (4), we can compute

$$f(2\Delta + t) = f(0) + \frac{t}{\Delta} Df(2\Delta) + \frac{t^2}{2\Delta^2} D^2f(2\Delta) \quad (7)$$

which is equivalent to the Taylor series

$$f(t) = x_0 + v_0 + \frac{1}{2}a_0t^2 \quad (8)$$

where x_0 , v_0 , and a_0 are parameters fit by the model. Any extra parameters, for example, for a 3rd order polynomial fit are zero, and thus omitted.

ADD CODE & FIGURE

1.3.2 Example: multi-dimensional trajectory

Here we consider a multi-dimensional example. The latent rule is at time $t_n \in \mathbb{N}$, we draw a square, where the top-left vertex is at the origin, with a side length of $s = t_n + 1$.

Treating the observation mechanism (*e.g.* the visual cortex) as an oracle, we identify eight features: the four coordinate-pairs which note the vertices of the square. One vertex (top left) never changes position – it is fixed at the origin. The top right vertex proceeds according to the rule $f_{tr}(t_n) = (0, t_n + 1)$. The bottom left vertex proceeds according to $f_{bl}(t_n) = (-1 - t, 0)$. Correspondingly, the bottom right vertex evolves according to $f_{br}(t_n) = (-1 - t, -1 - t)$.

Thus, each coordinate pair evolves either linearly or doesn't change (is constant). Finite difference approximations can be taken for each coordinate independently, and the rule can be determined after two steps.

ADD CODE & FIGURE

1.4 Representing a trajectory through function transformations

The idea of mapping vectors across a time-coordinate can be expanded to a larger space of problems. In this generalization, the latent rule involves defining a series of functions ordered in time: $[f_n(\mathcal{S})]_{n \in \mathbb{N}}$.

Each function acts on a surface \mathcal{S} , the specifics of which should not matter for the general formalism.

In the field of functional calculus, arithmetic can be defined on functions. The “sum” of two functions is the convolution $*$, and the “difference” of two functions is the cross-correlation \star .

If the rule can be modeled as a first-order transformation in function space, the evolution rule is:

$$f_{n+1}(t) = (f_n(t) \star f_{n-1}(t)) * f_n(t) \quad (9)$$

This formalism is an extension of the vector representation formalism (1.3), since convolution and cross-correlation are defined for all well-behaved functions.

1.4.1 Example: rate coding in biological neurons

Consider a biological neuron, where the firing rate is represented by $r(t)$, which receives an input $f(t)$ which is related to the firing rate by the following differential equation:

$$\frac{dr}{dt} = -kr + f(t) \quad (10)$$

where k is a parameter which describes how quickly the firing rate decays. Solving for $r(t)$ yields:

$$r(t) = F(t) \equiv \int_0^\infty f(t)e^{-kt} dt \quad (11)$$

which is the Laplace transform of $f(t)$. This formulation, in which biological neurons record temporal information in their rate of firing rate decay, has broad implications for trajectory prediction, since convolution and cross-correlation in the time domain are multiplication in the Laplace domain:

$$F_{n+1}(t) = F_n(t) \cdot \bar{F}_{n-1}(t) \cdot F_n(t) \quad (12)$$

2 Related Works

3 Methods

4 Results

5 Discussion

6 Conclusion